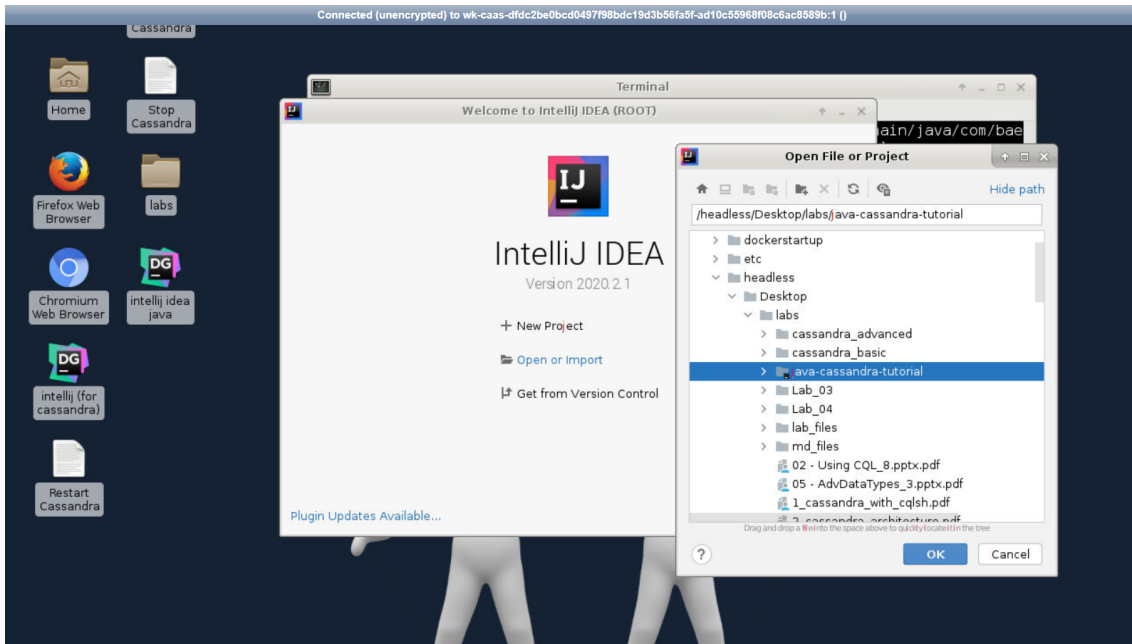


# Cassandra – Getting Started with Java

Cassandra is a great tool for storing time series data and I happen to be using it on my current project for that exact purpose.

There are several ways to use Cassandra from Java and many ways to improve performance, but here I just want to provide a simple “Getting Started” example. So here it is!

Code for this lab is located in `/headless/Desktop/labs/java-cassandra-tutorial`. Open project in IntelliJ IDE as shown below:



## Maven Project

Create a Java project, if using Maven, you can add the following dependencies to your pom.xml file:

```
<dependency>
  <groupId>com.datastax.cassandra</groupId>
  <artifactId>cassandra-driver-core</artifactId>
  <version>3.8.0</version>
</dependency>
```

Here is a simple Java example showing how to connect to Cassandra, create a keyspace, create a table, insert a row, and select a row:

```
import com.datastax.driver.core.*;

import java.time.Instant;
import java.time.ZoneId;
import java.util.Date;
import java.util.UUID;

public class CassandraV3Tutorial {
```

```

private final static String KEYSPACE_NAME = "example_keyspace";
private final static String REPLICATION_STRATEGY = "SimpleStrategy";
private final static int REPLICATION_FACTOR = 1;
private final static String TABLE_NAME = "example_table";

public static void main(String[] args) {

    // Setup a cluster to your local instance of Cassandra
    Cluster cluster = Cluster.builder()
        .addContactPoint("localhost")
        .withPort(9042)
        .build();

    // Create a session to communicate with Cassandra
    Session session = cluster.connect();

    // Create a new Keyspace (database) in Cassandra
    String createKeyspace = String.format(
        "CREATE KEYSPACE IF NOT EXISTS %s WITH replication = " +
            "'{'class':'%s','replication_factor':%s}";",
        KEYSPACE_NAME,
        REPLICATION_STRATEGY,
        REPLICATION_FACTOR
    );
    session.execute(createKeyspace);

    // Create a new table in our Keyspace
    String createTable = String.format(
        "CREATE TABLE IF NOT EXISTS %s.%s " + " " +
            "(id uuid, timestamp timestamp, value double, " +
            "PRIMARY KEY (id, timestamp)) " +
            "WITH CLUSTERING ORDER BY (timestamp ASC);",
        KEYSPACE_NAME,
        TABLE_NAME
    );
    session.execute(createTable);

    // Create an insert statement to add a new item to our table
    PreparedStatement insertPrepared = session.prepare(String.format(
        "INSERT INTO %s.%s (id, timestamp, value) values (?, ?, ?)",
        KEYSPACE_NAME,
        TABLE_NAME
    ));

    // Some example data to insert
    UUID id = UUID.fromString("1e4d26ed-922a-4bd2-85cb-6357b202eda8");
    Date timestamp = Date.from(Instant.parse("2018-01-01T01:01:01.000Z"));
    double value = 123.45;

    // Bind the data to the insert statement and execute it
    BoundStatement insertBound = insertPrepared.bind(id, timestamp, value);

```

```

session.execute(insertBound);

// Create a select statement to retrieve the item we just inserted
PreparedStatement selectPrepared = session.prepare(String.format(
    "SELECT id, timestamp, value FROM %s.%s WHERE id = ?",
    KEYSPACE_NAME,
    TABLE_NAME));

// Bind the id to the select statement and execute it
BoundStatement selectBound = selectPrepared.bind(id);
ResultSet resultSet = session.execute(selectBound);

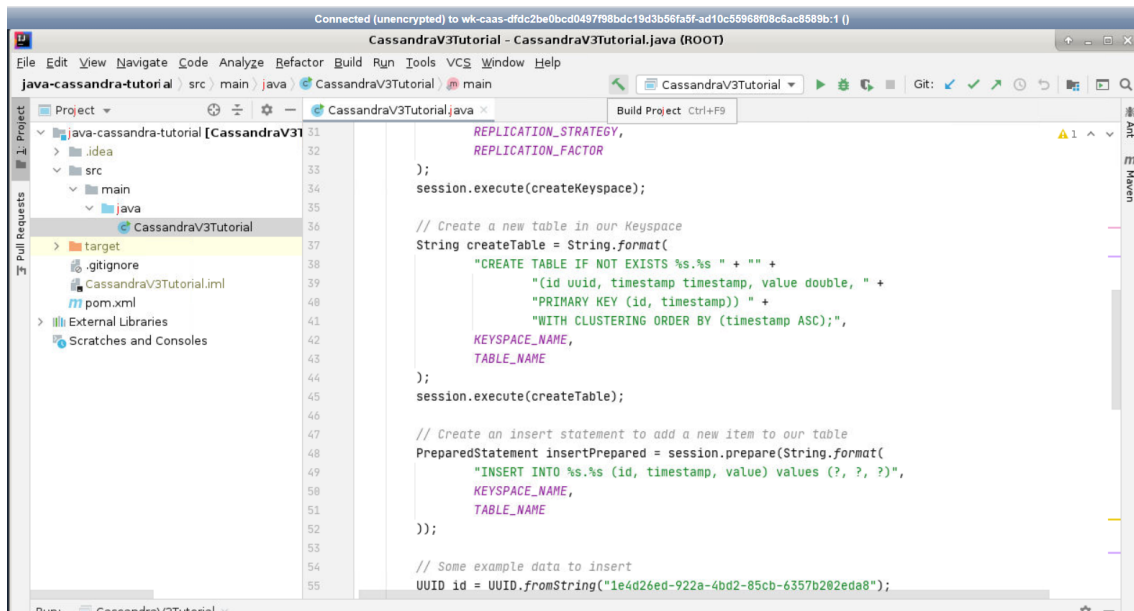
// Print the retrieved data
resultSet.forEach(row -> System.out.println(
    String.format("Id: %s, Timestamp: %s, Value: %s",
        row.getUUID("id"),
        row.getTimestamp("timestamp").toInstant().atZone(ZoneId.of("UTC")),
        row.getDouble("value"))));

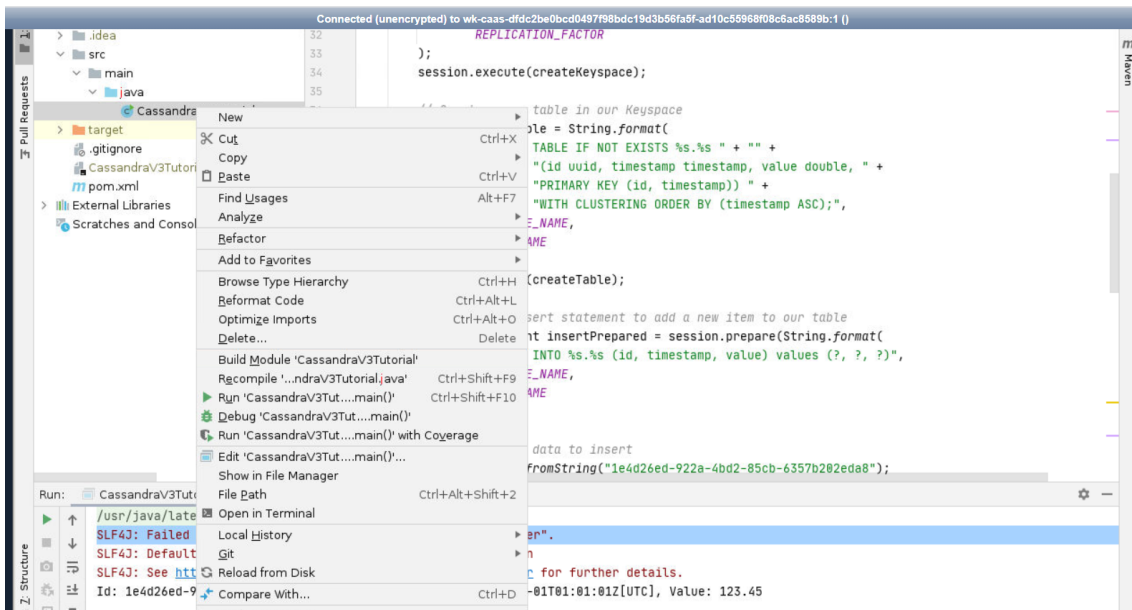
// Close session and disconnect from cluster
session.close();
cluster.close();
}
}

```

## Run Java Application

Make sure cassandra is running before running the example.





If you would like to look at the data in your local Cassandra database, you can use the [CQLSH](#) command line tool.

Type `cqlsh` in the terminal.

This will take you to a “cqlsh>” prompt:

```
Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 3.11.5 | CQL spec 3.4.4 | Native protocol v4]
Use HELP for help.
cqlsh>
```

To view all available Keyspaces:

```
DESCRIBE KEYSPACES;
```

You will now see our “example\_keyspace” in the list:

```
cqlsh> DESCRIBE KEYSPACES;

system_schema  system      system_traces
system_auth    system_distributed  example_keyspace
```

To switch to that Keyspace: `USE example_keyspace;`

```
cqlsh> USE example_keyspace;
cqlsh:example_keyspace>
```

To show all tables in the keyspace: `DESCRIBE TABLES;`

You should see the new table “example\_table”:

```
cqlsh:example_keyspace> DESCRIBE TABLES;

example_table
```

Now from the command line you can view the data in the table by using a select statement:

```
SELECT * FROM example_table;
```

Which will show the following information:

id	timestamp	value
1e4d26ed-922a-4bd2-85cb-6357b202eda8	2018-01-01 01:01:01.000000+0000	123.45

