Lab: Using Cassandra with cqlsh

To start working with Cassandra, let's start the **Cassandra Query Language (CQL)** shell . The shell interface will allow us to execute CQL commands to define, query, and modify our data. As this is a new cluster and we have turned on authentication and authorization, we will use the default `cassandra` and `cassandra` username and password, as follows:

```
  cqlsh localhost

Connected to PermanentWaves at localhost:9042.
[cqlsh 5.0.1 | Cassandra 3.11.2 | CQL spec 3.4.4 | Native protocol v4]
Use HELP for help.
cassandra@cqlsh>
```

First, let's tighten up security. Let's start by creating a new superuser to work with.

New users can only be created if authentication and authorization are properly set in the `cassandra.yaml` file:

```
  cassandra@cqlsh> CREATE ROLE cassdba WITH PASSWORD='flynnLives' AND LOGIN=true and SUPERUSER=true;
```

Now, set the default `cassandra` user to something long and indecipherable. You shouldn't need to use it ever again:

```
  cassandra@cqlsh> ALTER ROLE cassandra WITH PASSWORD='dsfawesomethingdfhdfshdlongandindecipherabledfhdfh';
```

Then, exit cqlsh using the `exit` command and log back in as the new `cassdba` user:

```
  cassandra@cqlsh> exit

cqlsh localhost -u cassdba -p flynnLives

Connected to PermanentWaves at localhost:9042.
[cqlsh 5.0.1 | Cassandra 3.11.2 | CQL spec 3.4.4 | Native protocol v4]
Use HELP for help.
cassdba@cqlsh>
```

Now, let's create a new keyspace where we can put our tables, as follows:

```
  cassdba@cqlsh> CREATE KEYSPACE fenagocassandra WITH replication =
    {'class':'SimpleStrategy', 'replication_factor':1}
  AND durable_writes = true;
```

Note

For those of you who have used Cassandra before, you might be tempted to build your local keyspaces with `SimpleStrategy` . `SimpleStrategy` has no benefits over `NetworkTopologyStrategy` , and is limited in that it cannot be used in a plural data center environment. Therefore, it is a good idea to get used to using it on your local instance as well.

With the newly created keyspace, let's go ahead and `use` it:

```
  cassdba@cqlsh> use fenagocassandra;
  cassdba@cqlsh:fenagocassandra>
```

Note

The cqlsh prompt changes depending on the user and keyspace currently being used.

Now, let's assume that we have a requirement to build a table for video game scores. We will want to keep track of the player by their `name` , as well as their `score` and `game` on which they achieved it. A table to store this data would look something like this:

```
  CREATE TABLE hi_scores (name TEXT, game TEXT, score BIGINT,
    PRIMARY KEY (name,game));
```

Next, we will `INSERT` data into the table, which will help us understand some of Cassandra's behaviors:

```
INSERT INTO hi_scores (name, game, score) VALUES ('Dad','Pacman',182330);
INSERT INTO hi_scores (name, game, score) VALUES ('Dad','Burgertime',222000);
INSERT INTO hi_scores (name, game, score) VALUES ('Dad','Frogger',15690);
INSERT INTO hi_scores (name, game, score) VALUES ('Dad','Joust',48150);
INSERT INTO hi_scores (name, game, score) VALUES ('Connor','Pacman',182330);
INSERT INTO hi_scores (name, game, score) VALUES ('Connor','Monkey Kong',15800);
INSERT INTO hi_scores (name, game, score) VALUES ('Connor','Frogger',4220);
INSERT INTO hi_scores (name, game, score) VALUES('Connor','Joust',48850);
INSERT INTO hi_scores (name, game, score) VALUES ('Avery','Galaga',28880);
INSERT INTO hi_scores (name, game, score) VALUES ('Avery','Burgertime',1200);
INSERT INTO hi_scores (name, game, score) VALUES ('Avery','Frogger',1100);
INSERT INTO hi_scores (name, game, score) VALUES ('Avery','Joust',19520);
```

Now, let's execute a CQL query to retrieve the scores of the player named `Connor` :

```
cassdba@cqlsh:fenagocassandra> SELECT * FROM hi_scores WHERE name='Connor';
 name   | game        | score
--------+-------------+--------
 Connor |     Frogger | 4220
 Connor |       Joust | 48850
 Connor | Monkey Kong |  15800
 Connor |      Pacman | 182330
(4 rows)
```

That works pretty well. But what if we want to see how all of the players did while playing the `Joust` game, as follows:

```
cassdba@cqlsh:fenagocassandra> SELECT * FROM hi_scores WHERE game='Joust';

InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot execute
this query as it might involve data filtering and thus may have unpredictable performance.
If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING"
```

Note

As stated in the preceding error message, this query could be solved by adding the `ALLOW FILTERING` directive. Queries using `ALLOW FILTERING` are notorious for performing poorly, so it is a good idea to build your data model so that you do not use it.

Evidently, Cassandra has some problems with that query. We'll discuss more about why that is the case later on. But, for now, let's build a table that specifically supports querying high scores by `game` :

```
CREATE TABLE hi_scores_by_game (name TEXT, game TEXT, score BIGINT,
  PRIMARY KEY (game,score)) WITH CLUSTERING ORDER BY (score DESC);
```

Now, we will duplicate our data into our new query table:

```
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Dad','Pacman',182330);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Dad','Burgertime',222000);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Dad','Frogger',15690);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Dad','Joust',48150);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Connor','Pacman',182330);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Connor','Monkey Kong',15800);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Connor','Frogger',4220);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Connor','Joust',48850);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Avery','Galaga',28880);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Avery','Burgertime',1200);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Avery','Frogger',1100);
INSERT INTO hi_scores_by_game (name, game, score) VALUES ('Avery','Joust',19520);
```

Now, let's try to query while filtering on `game` with our new table:

```
cassdba@cqlsh:fenagocassandra> SELECT * FROM hi_scores_by_game
                         WHERE game='Joust';

 game  | score | name
-------+-------+--------
 Joust | 48850 | Connor
 Joust | 48150 |    Dad
 Joust | 19520 |  Avery

(3 rows)
```

As mentioned previously, the following chapters will discuss why and when Cassandra only allows certain `PRIMARY KEY` components to be used in the `WHERE` clause. The important thing to remember at this point is that in Cassandra, tables and data structures should be modeled according to the queries that they are intended to serve.